



בית. תוכנה. חברה.

[info@zion-net.co.il](mailto:info@zion-net.co.il)



# Generative AI in C#

## Harnessing Large Language Models for Enhanced Development

# Is Generative AI the End for Developers?





# About Me

## ▶ Alon Fliess:

- ▶ CTO of ZioNet
- ▶ More than 30 years of hands-on experience
- ▶ Microsoft Regional Director & Microsoft Azure MVP



# About ZioNet

- ▶ A Professional Software service provider company
- ▶ The home for hi-potential novice developers and expert leaders
  - ▶ We support our developers' growth, provide them with professional and personal mentoring
- ▶ ZioNet management has over 20 years of experience
  - ▶ We strive to fulfill the need by ensuring developers have the best first-job experience!

# Introduction to AI





# Overview of Different Types of AI

- ▶ **Algorithm-based AI:** Uses rule-based decision-making systems
- ▶ **Supervised Learning:** Learns from a labeled dataset
- ▶ **Unsupervised Learning:** Finds hidden patterns in data
- ▶ **Reinforcement Learning:** Improves via reward-based feedback
- ▶ **Hybrid AI:** Combines different AI methodologies
- ▶ **Large Language Models:** Generates text by learning from a tremendous amount of text

# What is a Neural Network?

- ▶ Inspired by the human or animal brain
  - ▶ Far from being the same



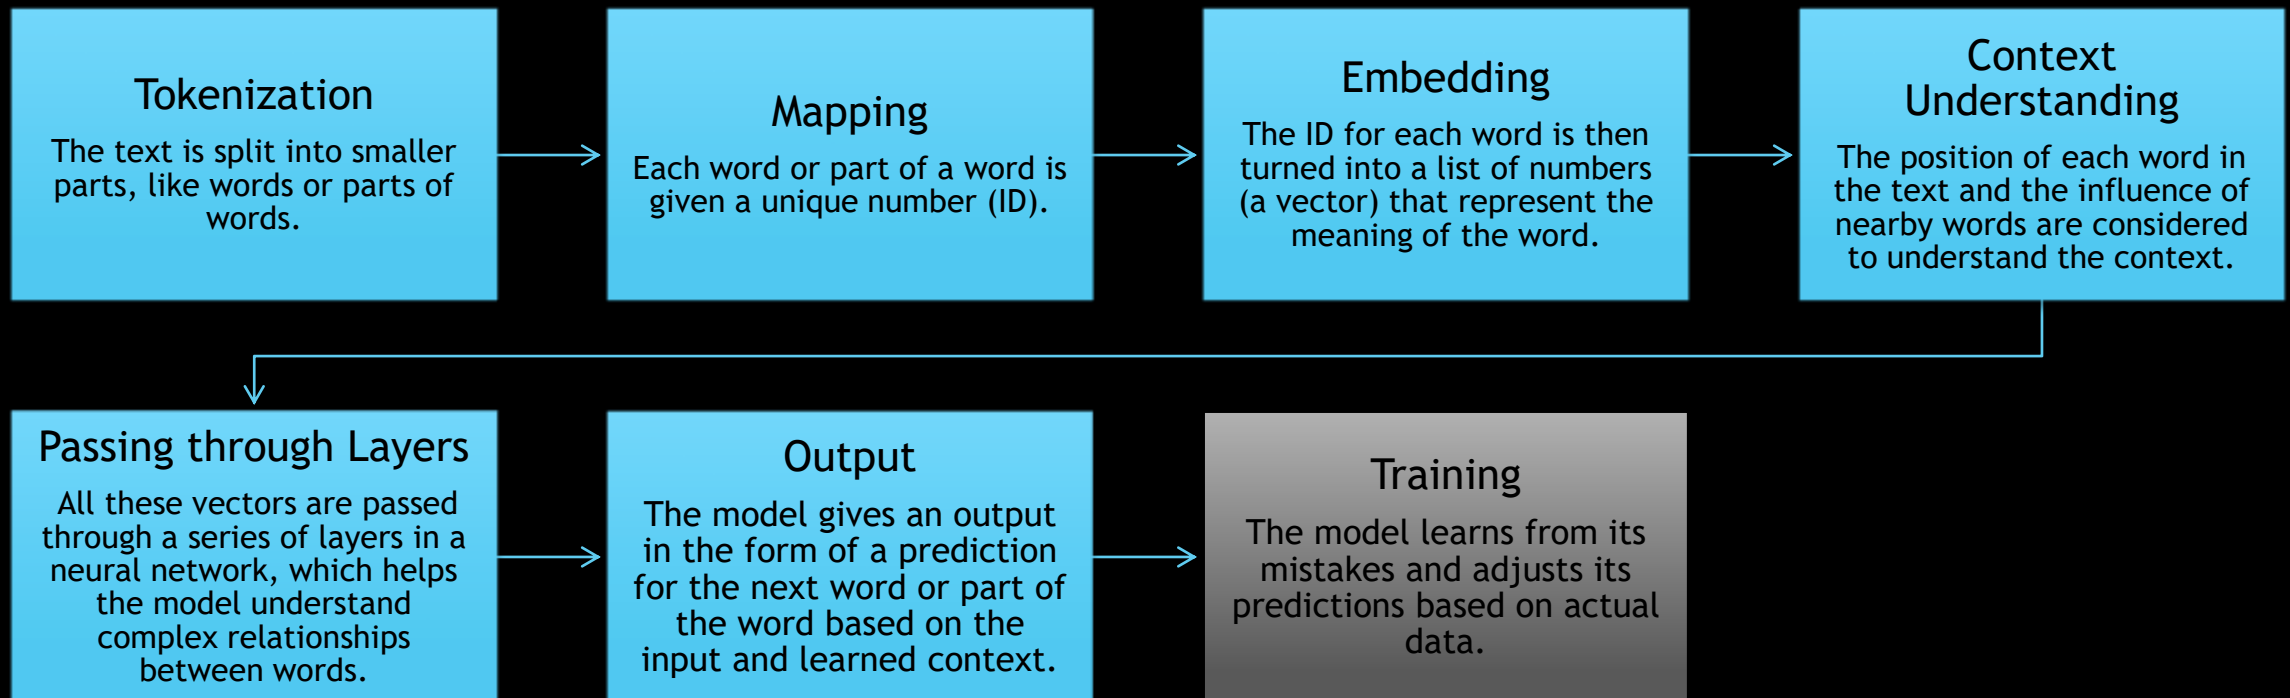


# Large Language Model Overview ([GPT](#), [LLaMA](#), [LaMDA](#), [PaLM](#))

- ▶ Capture the semantics of a language
- ▶ Trained with very large amount of data
- ▶ All you need is Attention... and position (context)
- ▶ The model predicts the next word using probability, based on the input text
- ▶ The next word (token) is predicted using the original input and all the word that were generated before



# LLM Processing (GPT)



# Tokenization - Explained

Tokens	Characters
34	182

Tokenization is the process of demarcating and possibly classifying sections of a string of input characters. The resulting tokens are then passed on to some other form of processing

Tokens	Characters
34	182

[30642, 1634, 318, 262, 1429, 286, 1357, 5605, 803, 290, 5457, 1398, 4035, 9004, 286, 257, 4731, 286, 5128, 3435, 13, 383, 7186, 16326, 389, 788, 3804, 319, 284, 617, 584, 1296, 286, 7587]



# What are the Base Models

([OpenAI](#), [Azure](#))

- ▶ Some predefined models were trained to do specific tasks
- ▶ There are variations:
  - ▶ The size (input + output) tokens
  - ▶ The speed
  - ▶ The usage price
  - ▶ The ability to be fine-tuned

# Can I have my Fined-Tuned Models?

## ▶ Why Fine-Tune?

- ▶ Improve the model performance on specific tasks
- ▶ Adapt the model to new data
- ▶ Customize the model's behavior

## ▶ When to Fine-Tune?

- ▶ When the Pre-Trained Model is not performing well, or the grounding (system) prompt is too large
- ▶ When you have a specific task or data

## ▶ How to Fine-Tune?

- ▶ Use tools like Azure Machine Learning Studio or OpenAI's Python client for fine-tuning
- ▶ Specify the Base Model and Dataset
- ▶ Monitor Fine-Tuning Job and retrain
- ▶ Use Fine-Tuned Model for Predictions

# Learn how to Generate or Manipulate Text

- ▶ **Classification**, such as sentiment
- ▶ **Generation** - Create a text or formatted text (JSON, XML)
- ▶ **Conversation** - Chat to get information, give commands, or generate and manipulate the result
- ▶ **Transformation**:
  - ▶ Language Translation
  - ▶ Text to emoji
  - ▶ Any format to any format
- ▶ **Summarization** - reduce the size of text
- ▶ **Completion** - complete a statement
- ▶ **Code** - Use Codex to generate, complete, or manipulate source code





LLM for Code  
is like SoC for  
Hardware


# Introduction to Prompt Engineering

- ▶ What is the **first** thing that comes to your mind when I say “<prompt>?”
- ▶ Prompt Engineering is the art of crafting inputs to get desired outputs from AI models
- ▶ It's a crucial part of using AI models effectively
  - ▶ The design of the prompt can greatly influence the model's response.
- ▶ Examples:
  - ▶ If you want a list, start your prompt with a numbered list.
  - ▶ If you want a specific format, provide an example in that format.
- ▶ It often involves a lot of trial and error
  - ▶ Different prompt strategies may work better for different tasks
- ▶ Less is more!

# Prompt Engineering Recommendations

- ▶ **Goal:** Define what you want from the model
- ▶ **Instructions:** Be clear and explicit
- ▶ **Examples:** Use them for specific formats or styles
- ▶ **Iterate:** Experiment with different prompts
- ▶ **Guidance:** Use system-level and user-level instructions
- ▶ **Settings:** Adjust temperature and max tokens as needed

# Playground - ChatGPT and Azure



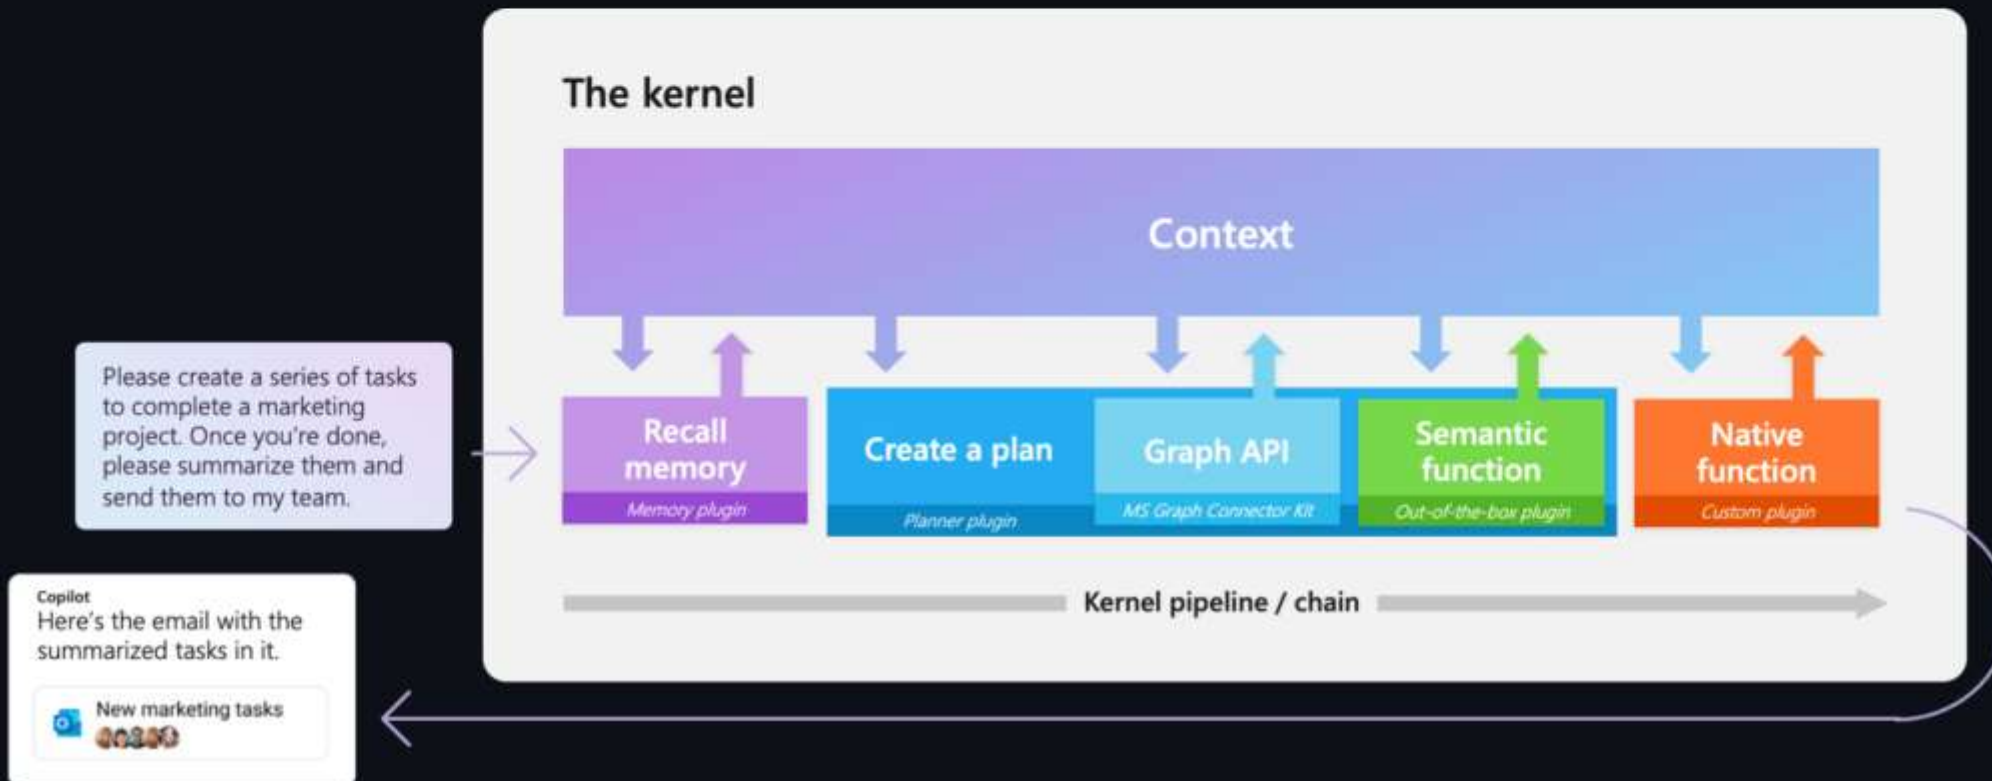
- ▶ To get started, use [ChatGPT](#) or [Azure Open AI playground](#) - Demo
- ▶ You can generate your boilerplate code for grounding and examples
- ▶ You can play with fine-tuning parameters

[pypi v0.4.7.dev0](#)
[Nuget package](#)
[dotnet-ci-docker passing](#)
[dotnet-ci-windows passing](#)
[license MIT](#)
[Discord 380 online](#)

[Semantic Kernel](#) is an SDK that integrates Large Language Models (LLMs) like [OpenAI](#), [Azure OpenAI](#), and [Hugging Face](#) with conventional programming languages like C#, Python, and Java. Semantic Kernel achieves this by allowing you to define [plugins](#) that can be chained together in just a [few lines of code](#).

What makes Semantic Kernel *special*, however, is its ability to *automatically* orchestrate plugins with AI. With Semantic Kernel [planners](#), you can ask an LLM to generate a plan that achieves a user's unique goal. Afterwards, Semantic Kernel will execute the plan for the user.

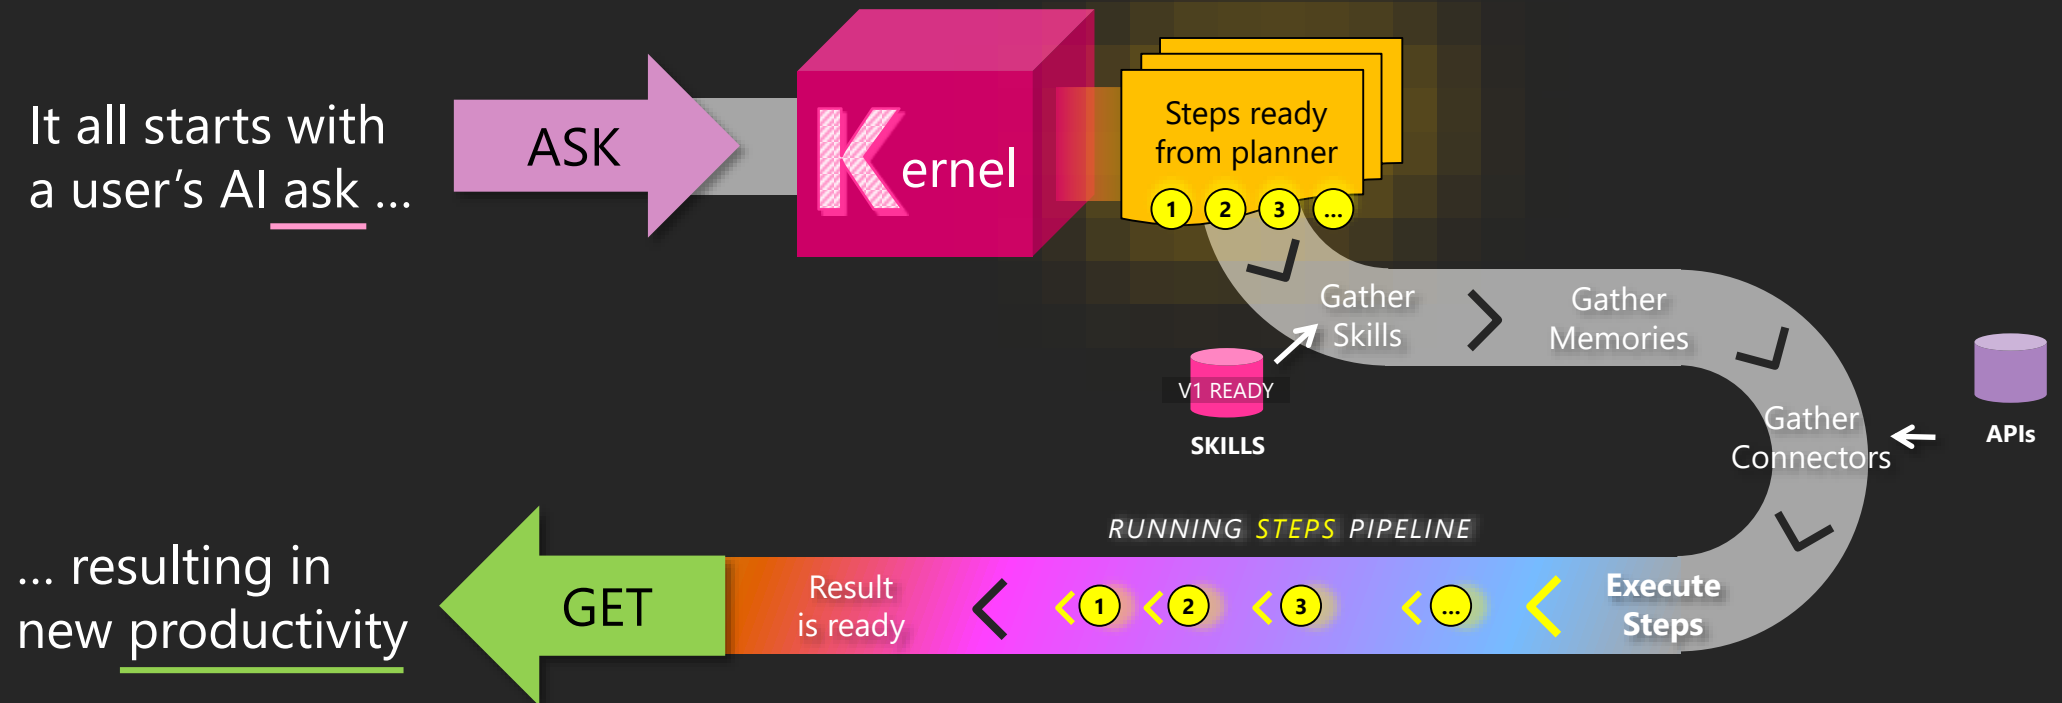
Please star the repo to show your support for this project!





# Goals-First AI


It all starts with  
a user's AI ask ...



# Semantic Kernel Main Concepts

- ▶ **Prompt Functions:**
  - ▶ Define interaction patterns with LLMs
- ▶ **Native Functions:**
  - ▶ Enable direct code execution by AI
- ▶ **Plugins:**
  - ▶ Custom-built, modular elements enabling specialized LLM task handling
- ▶ **Memory:**
  - ▶ Contextual data repository; supports key-value pairs and semantic embeddings
- ▶ **Connectors:**
  - ▶ Interface with external data and APIs
- ▶ **Planners:**
  - ▶ Orchestrate tasks, manage complex tasks through intelligent LLM planning
- ▶ **Agents:**
  - ▶ Autonomous entities executing orchestrated tasks

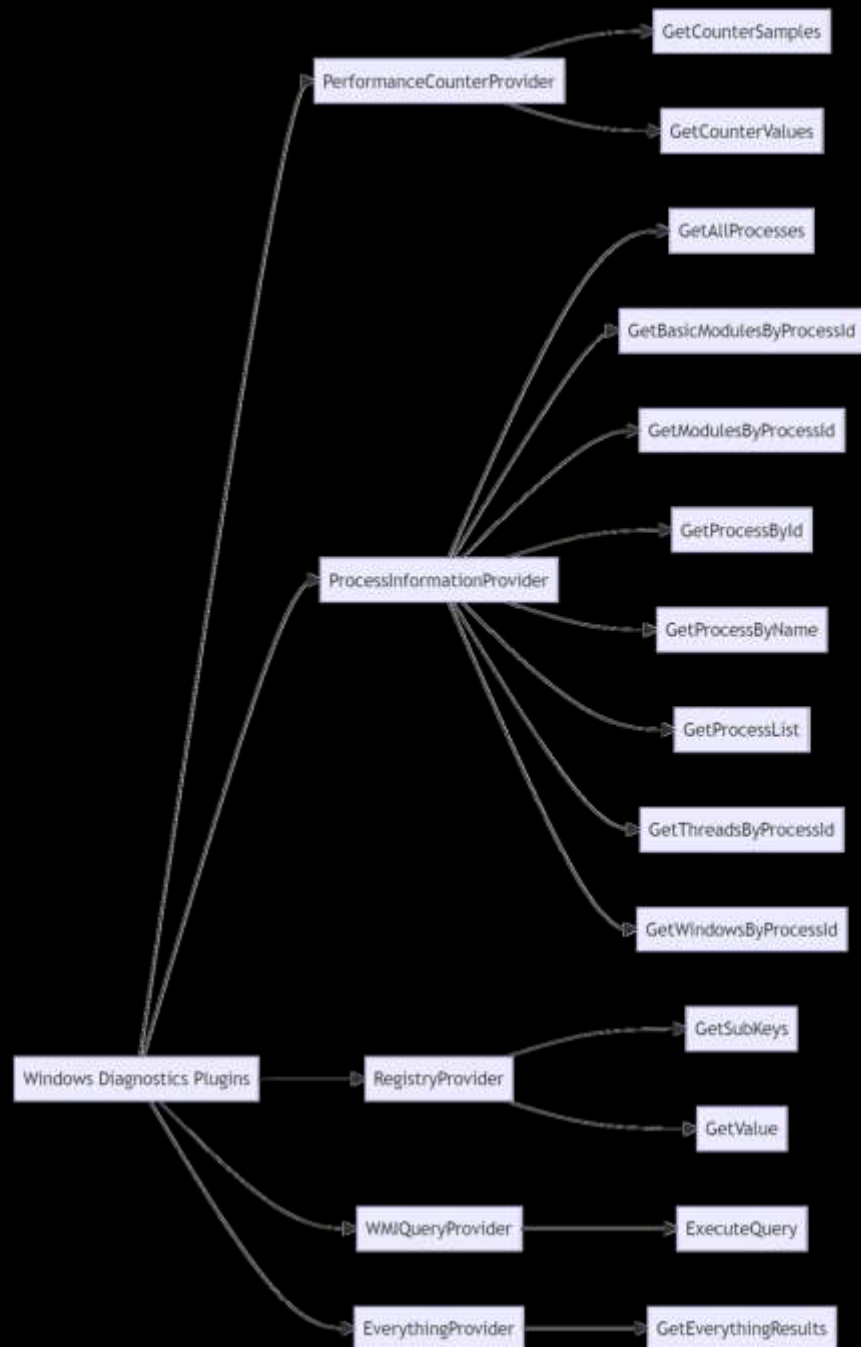
# Demo



# Windows Troubleshooting Plugin Capabilities

- ▶ **Event Log:** Query Windows Event Logs
- ▶ **File System:** Read and search files
- ▶ **Registry:** Access and query keys
- ▶ **Performance:** Monitor system metrics
- ▶ **Network Info:** Retrieve connection details
- ▶ **Process Info:** List running processes
- ▶ **Service Info:** Query system services
- ▶ **WMI Query:** Execute WMI commands
- ▶ **Everything:** Advanced file search
- ▶ **Paging:** Supports result pagination
- ▶ **Tenant/PC ID:** Requires GUIDs.
- ▶ **Multi-Function:** Multiple queries in one call
- ▶ **Parameter Flex:** Customizable parameters
- ▶ **Truncation:** Handles truncated results
- ▶ **Max Token:** Customizable data size.

# The Windows Troubleshooting Plugin





# Lesson Learned

- ▶ The first project iteration: An OpenAI ChatGPT plugin
  - ▶ Before Semantic Kernel
  - ▶ Had to bridge LLM to a C# Code
    - ▶ Simple APIs, Json, Reflection, Default Values
    - ▶ Needed to reduce the description load due to token limitation
    - ▶ No Planner
- ▶ The second project iteration: Use multiple LLM and SK
  - ▶ Work in progress
  - ▶ Requires lots of fine tuning - debug while you go
  - ▶ Expensive
  - ▶ For agents → Use Asynchronous Model or a batch processing

# Lessons Learned from Developing Plugin

## ▶ *Plugin API Design*

- **Be Systematic:** Stick to one or a few APIs for consistency
- **OpenAPI Description:** Provide a comprehensive Open API specification

## ▶ *Plugin Manifest*

- **Examples:** Include examples for each query method
- **Instructions:** Update examples and instructions if ChatGPT calls with the wrong data schema

## ▶ *Data Handling*

- **Paging:** Implement paging capabilities
- **Truncation:** Use HTTP code 206 and a special message to indicate truncated results
- **Important Data:** Always return key data like I do in the Windows Troubleshooting with Tenant Id and PC ID

## ▶ *Performance & Limitations*

- **Trial and Error:** Extensive testing is crucial
- **Size Limits:** Be aware of total size and per JSON element limits
- **Resource Management:** Be cautious of overusing ChatGPT 4 resources

## ▶ *Dynamic Content*

- For complex plugins, dynamically generate the Open API specification and plugin manifest

# Lessons Learned from Embedding LLM into Applications

## ▶ **Grounding & Serialization**

- Provide accurate grounding and a JSON schema describing the result
- Use Semantic Kernel Prompt and Native function – use string and Json parameters

## ▶ **Model Selection**

- Use GPT3/3.5/4 for chat functionalities
- Use other models for embedding, image creation and recognition

## ▶ **Performance & Cost**

- GPT4 is more accurate but costly and slower
- Consider fine-tuned models if applicable – high cost for a good result

## ▶ **Message Handling**

- Handle truncated messages with a continuation strategy – less important on 120K tokens
- Manage message size by counting tokens and removing history
- Use summary messages to replace original history if needed
  - Use Memory (Embedding/Vectorization)

# Overcome the Model Limitation

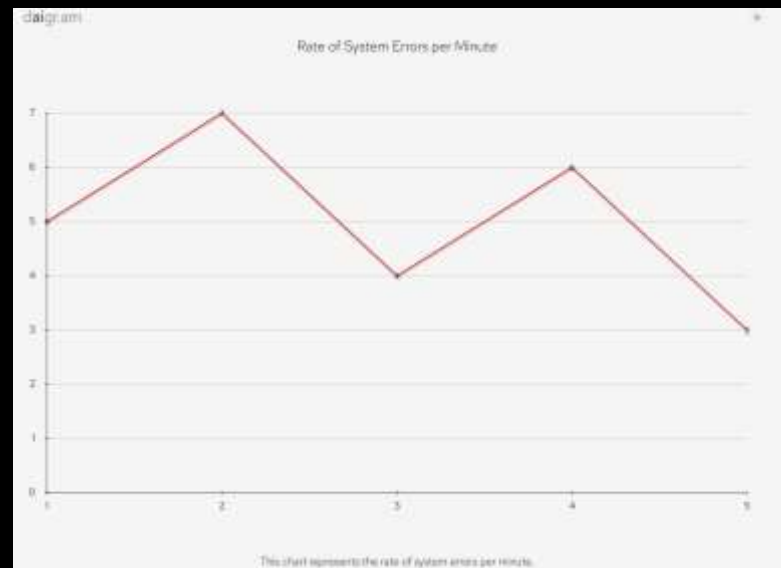
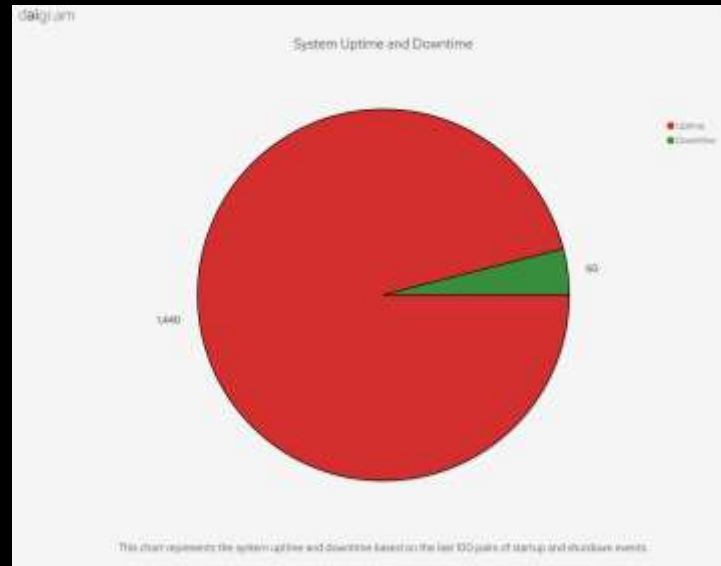
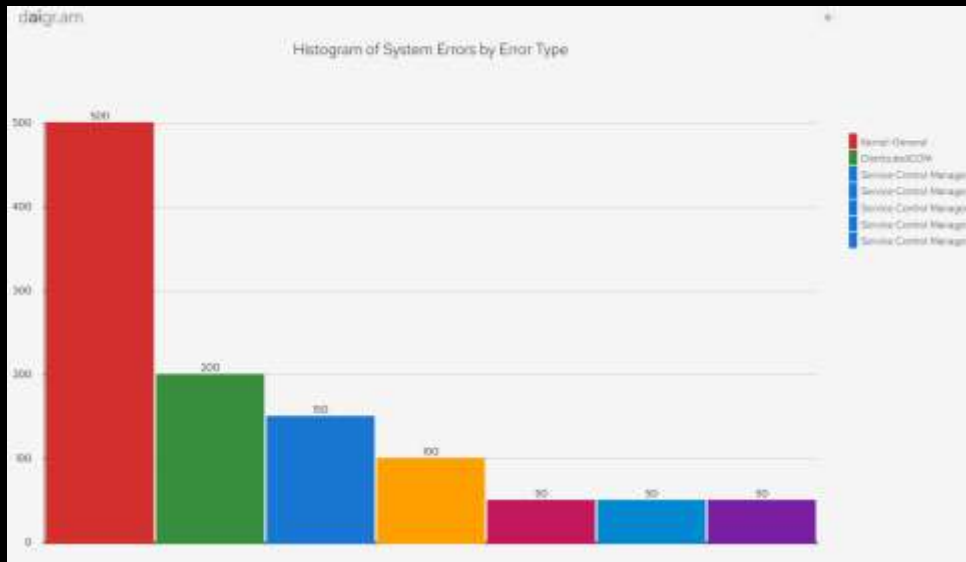
- ▶ Token Limitation
  - ▶ Use Summarization
  - ▶ Use Memory (store important facts)
  - ▶ Use Retrieval Augmented Generation (RAG) + Vector and other search
- ▶ Overcome the lack of current information
  - ▶ Use Bing, or other search Plugins
- ▶ Overcome cost and availability
  - ▶ Use GPT3.5 Turbo
  - ▶ Use Open-Source Models
  - ▶ Prompt engineer cheap models (3.5) with GPT 4

# The Windows Event Log Plugin

- ▶ <https://github.com/alonf/WindowsEventLogChatGPTPlugin.git>
- ▶ Retrieve specific events from the Windows Event Log using XPath queries.
- ▶ Supports all major log names: Application, Security, Setup, System, and ForwardedEvents.
- ▶ Use it to solve problems and get information about your Windows system status.
- ▶ The plugin supports paging. It estimates the number of tokens and limits the result.



# Examples



# Developing ChatGPT Plugin Using C#

- ▶ Use ASP.NET Minimal or Controller based API
- ▶ Use [YamlDotNet](#) to convert Json Open API specification to yaml based
  - ▶ The Open API specification, the plugin manifest and the icon file can come from the file system, or as a HTTP query

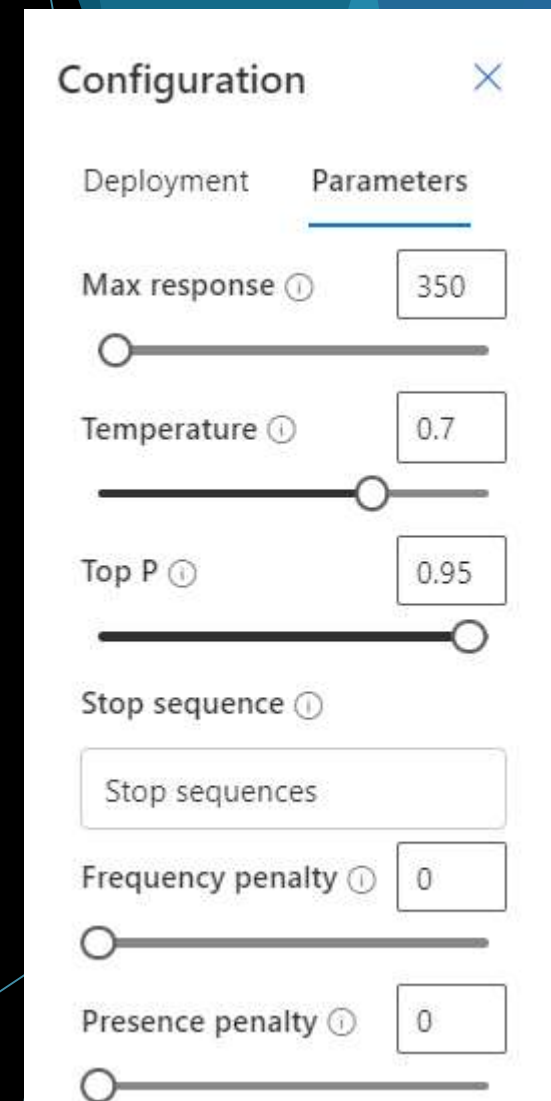
```
app.UseStaticFiles(new StaticFileOptions
{
    FileProvider = new PhysicalFileProvider(
        Path.Combine(app.Environment.WebRootPath, "OpenAPI")),
    RequestPath = "/OpenAPI"
});
```

# Developing ChatGPT Plugin Using C#

- ▶ Route HTTP Request to a function
  - ▶ For a simple plugin, it is just a get request
  - ▶ For A complex plugin, use POST with a body and have your own route
- ▶ In the Windows Troubleshooting plugin, I use a map of providers and made the call to the specific function using reflection
  - ▶ I use validation to make sure the data size and type is correct
  - ▶ I provide extensive error message
- ▶ Use correlation id for local development (Tenant ID, PC ID)
- ▶ Use OAuth for released plugin

# Controlling the Output layer

- ▶ Max response: defines the token limit for the model's response
  - ▶ One token is roughly equivalent to 4 English characters
- ▶ Temperature: Controls the randomness of the model's responses.
  - ▶ Lower values result in more deterministic responses, higher values lead to creativity
- ▶ Top P: Another parameter to control randomness
  - ▶ Lower values make the model choose more likely tokens
- ▶ Stop sequence:
  - ▶ Specifies a sequence at which the model should stop generating a response
- ▶ Frequency penalty: Reduces the likelihood of the model repeating the same text
  - ▶ by penalizing tokens that have appeared frequently
- ▶ Presence penalty:
  - ▶ Encourages the model to introduce new topics in a response
    - ▶ by penalizing any token that has appeared in the text so far.



The screenshot shows the 'Configuration' panel of the OpenAI Playground. It has two tabs: 'Deployment' and 'Parameters', with 'Parameters' currently selected. The parameters are as follows:

Parameter	Value
Max response	350
Temperature	0.7
Top P	0.95
Stop sequence	Stop sequences
Frequency penalty	0
Presence penalty	0

Each parameter has a slider control below its input field. The 'Stop sequence' field is a text input box containing the text 'Stop sequences'.

# Cost, Privacy & Security

- ▶ ChatGPT and Azure OpenAI services can be used without donating your data
  - ▶ For the Public ChatGPT you can ask to opt-out
  - ▶ For the API, You can ask to opt-in
- ▶ ChatGPT 4 is a very hi-cost model (become cheaper)
  - ▶ You can use ChatGPT 4 to create prompt and examples for ChatGPT 3.5
- ▶ You can host your model on-premise, however:
  - ▶ The Open-Source models do not contain the latest ChatGPT 3.5 and 4
  - ▶ You may train your own model - costly

# Summary

- ▶ **AI Types**: Explored diverse AI forms
- ▶ **Neural Networks & LLMs**: Discussed their functionality
- ▶ **Base Models & Fine-Tuning**: Highlighted fine-tuning's importance
- ▶ **Prompt Engineering**: Introduced crafting effective prompts
- ▶ **Semantic Kernel**: Your LLM Swiss army tool
- ▶ **Application Transformation**: Extending applications with AI
- ▶ **Autonomous Agents**: The future of AI systems

LLM ↔ The Software System on a Chip





The End

